

# CSCI 379 Computer Networking


---

## Textbook Notes 6.1-6.3

---

### 6.1: Introduction to the Link Layer

- In this chapter, any device that runs a link-layer protocol will be referred to as a *node*
  - Nodes include hosts, routers, switches, and Wi-Fi access points
- The communication paths between nodes connecting them will be referred to as *links*
- Over any given link, a transmitting node encapsulates the datagram in a *link-layer frame* and transmits the frame into the link
- **6.1.1: Services Provided by the Link Layer**
  - Provided services can vary from one link-layer protocol to the next
  - Possible services provided by link-layer protocols include
    - *Framing*, which refers to the encapsulation of a network layer datagram in a link-layer frame prior to transmission
    - *Link access*, such as through the medium access control (MAC) protocol, which deals with managing access and coordination of frame transmission over different nodes
    - *Reliable Delivery*, which similar to TCP refers to the guarantees made by the link-layer in terms of guaranteed and in-order delivery using acknowledgments and retransmissions
    - *Error Detection and Correction*, which functions like the checksum properties in the network layer, but generally uses more sophisticated error checking and correcting methodologies
- **6.1.2: Where is the Link Layer Implemented**
  - For the most part, in user-side end systems, the link-layer is implemented in the form of a *network adapter* and *network interface card (NIC)*



- At its heart is the link-layer controller, which is generally a single, special-purpose chip that implements many of the link-layer's services
- So, as we can see, the majority of a link-layer controller's functionality is implemented in hardware
- Until the late 1990s, network adapters were almost exclusively physically separate cards (PCMCIA, PCI, PCIe), but in the modern day, motherboards almost always have an on-board LAN configuration
- Nowadays, many motherboards. even have on-board Wi-Fi configurations
- While the link-layer *is* largely implemented in hardware, software side implementations also assist in controller interrupts, error handling and datagram passing to the network layer

---

## 6.2: Error-Detection and -Correction Techniques

- We have noted throughout the text that error-detection and -correction techniques are two services that are often provided by the link layer, as well as the transport layer
- In this section, we will examine some of the simplest techniques that can be used to detect, and in some cases, correct bit errors

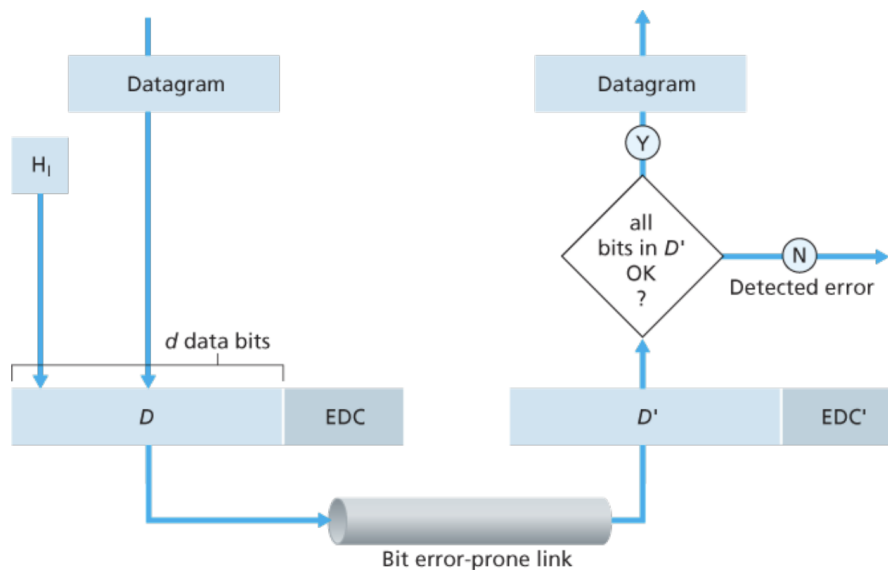



Figure 6.3 Error-detection and -correction scenario

- Here, we can see that at the sending node,  $D$  represents data to be protected against bit errors
  - This data will be augmented with error-detection and -correction bits (EDC)
- The data to be protected typically does not only include the datagram passed down from the network layer, but also link-level addressing information and other fields in the link frame header
- At the receiving node,  $D'$  and  $EDC'$  are received, which might differ from the original  $D$  and  $EDC$
- Now, the receiver's challenge is to determine whether the  $D'$  which it has received is the same as the original  $D$
- Error-detection and -correction techniques allow the receiver to sometimes, *but not always*, detect when bit errors have occurred
- There can still be *undetected bit errors*, which means the receiver might deliver a corrupted datagram to the network layer, or be unaware the contents of a field in the frame's header has been corrupted
- Thus, we want to choose an error-detection scheme that keeps the probability of such occurrences small
- In general, more sophisticated error-detection and -correction algorithms incur more overhead due to the need for more computation and transmission of EDC bits

- 6.2.1: Parity Checks

- The simplest form of error detection is the use of a single *parity bit*
- Suppose that the information we are sending,  $D$ , has  $d$  bits
  - In an *even parity scheme* the sender includes one additional bit such that the total number of 1s in the  $d + 1$  bits is even
  - For *odd parity schemes*, the parity bit value is chosen such that there is an odd number of 1s



**Figure 6.4 One-bit even parity**

- Receiver operation is also fairly simple when utilizing a parity bit
- The receiver need only count the number of 1s in the received  $d+1$  bits
  - If an odd number of 1-valued bits are found within an even parity scheme, the receiver knows that at least one bit error has occurred
  - More precisely, the receiver knows that some *odd* number of bit errors have occurred
- But what happens when an even number of bit errors occurs?
  - This would result in an undetected error
- If the probability of bit errors is small, and errors can be assumed to occur independently from one another, the probability of having multiple bit errors in a single packet is extremely small
- However, measurements have shown that, rather than occurring independently, errors are often clustered together in *bursts*
  - Under burst conditions, the probability of undetected errors in a frame protected by single-bit parity can approach 50%

- Clearly, a more robust error detection scheme is needed for practical use, but before examining such a scheme, we will consider a generalization of one bit parity that will provide us with some insight into error-correction techniques

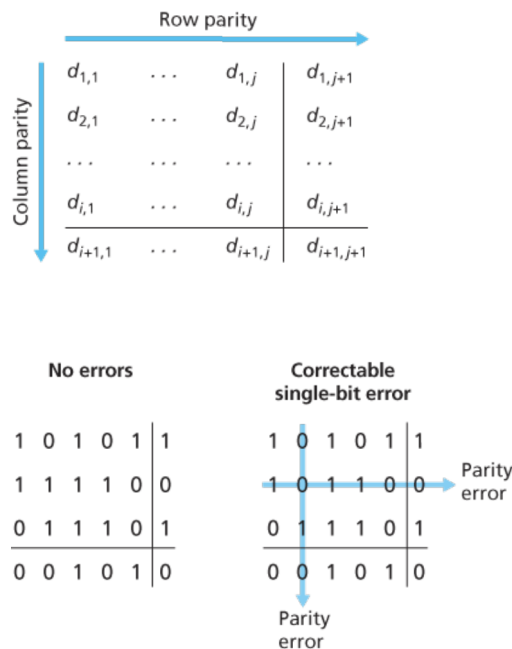


Figure 6.5 Two-dimensional even parity

- As shown in the above diagram, a two-dimensional generalization can help us better understand error correction techniques
- Here, the  $d$  bits in  $D$  are divided into  $i$  rows and  $j$  columns, and a parity value will be computed in each row and each column
  - The resulting  $i + j + 1$  parity bits comprise the link-layer frame's error detection bits
- Now, suppose that a single bit error occurs in the original  $d$  bits of information
- In this scheme, the parity of both the column and the row containing the flipped bit will be in error
  - Therefore, the receiver can use the indices of the row and column bits in error to not only detect that a bit is in error, but also identify the corrupted bit, and *correct* the error
- Two dimensional parity schemes can also detect, but not correct, any combination of two errors in a single packet

- The ability of the receiver to both detect and correct errors is known as *forward error correction (FEC)*
- FEC techniques are valuable since they can decrease the number of sender transmissions required
- It also avoids the round-trip propagation of ACK and NAK packets and retransmissions, which can prove integral to real-time web applications

- **6.2.2: Checksumming Methods**

- In *checksumming techniques*, the  $d$  bits of data are treated as a sequence of  $k$ -bit integers
- One simple checksumming method would be to sum these  $k$ -bit integers and use the resulting sum as the error-detection bits
  - The *Internet checksum* is based on this approach; bytes of data are treated as 16-bit integers and then summed
- Checksumming methods require relatively little overhead
  - For example, the checksums in TCP and UDP use only 16-bits
- However, they provide relatively weak protection against errors when compared to other error detection and correction methods that will be later discussed

- **6.2.3: Cyclic Redundancy Check (CRC)**

- A widely used error detection technique in modern computer networks is based on *cyclic redundancy check (CRC) codes*
  - These are also called *polynomial codes*, since it is possible to view the bit string to be sent as a polynomial whose coefficients are the 0 and 1 values in the bit string, with operations on the bit string interpreted as polynomial arithmetic
- **CRC codes operate as follows**
  - Consider the  $d$ -bit piece of data,  $D$ , that the sending node wants to send to the receiving node
  - The sender and receiver must first agree on an  $r + 1$  bit pattern, known as a *generator* which will be denoted by  $G$
  - The leftmost bit of  $G$  must be a 1

- For  $D$ , the sender will choose  $r$  additional bits,  $R$ , and append them to  $D$  such that the resulting  $d + r$  bit pattern is exactly divisible by  $G$  using modulo-2 arithmetic




Figure 6.6 CRC

- The process of error checking is fairly simple with CRCs
  - The receiver computes:
    - $$\frac{d + r}{G} \text{ mod } 2$$
    - If the result of this computation is nonzero, the receiver knows that an error has occurred, and otherwise the data is accepted as being correct
- Given  $D$  and  $R$ , the formula  $D \cdot 2^r \text{ XOR } R$  yields the  $d + r$  bit pattern shown in Figure 6.6
- How will the sender compute  $R$ ?
  - We want to find  $R$  such that there is an  $n$  which satisfies
    - $$\begin{aligned} D \cdot 2^r \oplus R &= nG \\ D \cdot 2^r &= nG \oplus R \end{aligned}$$
    - Thus, if we divide  $D \cdot 2^r$  by  $G$ , the value of the remainder is precisely  $R$
    - So, we can calculate  $R$  as
      - $$R = \text{remainder}\left(\frac{D \cdot 2^r}{G}\right)$$
  - Below is an example of the calculations described above
  -




Figure 6.7 A sample CRC calculation

- International standards have been defined for 8, 12, 16, and 32-bit generators
- The CRC-32 32-bit standard, which has been adopted in a number of link layer IEEE protocols, uses a generator of  $G_{CRC-32} = 100000100110000010001110110110111$
- Each of the CRC standards can detect burst errors of fewer than  $r + 1$  bits
- Each of the CRC standards can also detect any odd number of bit errors

### 6.3: Multiple Access Links and Protocols

- There are two different types of network links
  - A *point-to-point* link consists of a single sender at one end of the link and a single receiver at the other end of the link
    - Many link-layer protocols have been designed for point-to-point links, such as the *point-to-point protocol (PPP)* and *high-level data link control (HDLC)*
  - The second type of link, a *broadcast link*, can have multiple sending and receiving nodes which are all connected to the same, single, shared broadcast channel
    - Here, the term *broadcast* is used here because when any one node transmits a frame, each of the other nodes will receive a copy
    - Ethernet and wireless LANs are examples of broadcast link-layer technologies
- Now, we must consider how to coordinate the access of multiple sending and receiving nodes to a shared broadcast channel; the *multiple access problem*



- Television has been using the idea of broadcasting since long before the advent of computers
  - However, traditional television is a one-way broadcast link, while nodes on a computer network broadcast channel can both send *and* receive
  - Networks have protocols known as *multiple access protocols*, by which nodes regulate their transmission into the shared broadcast channel
- 

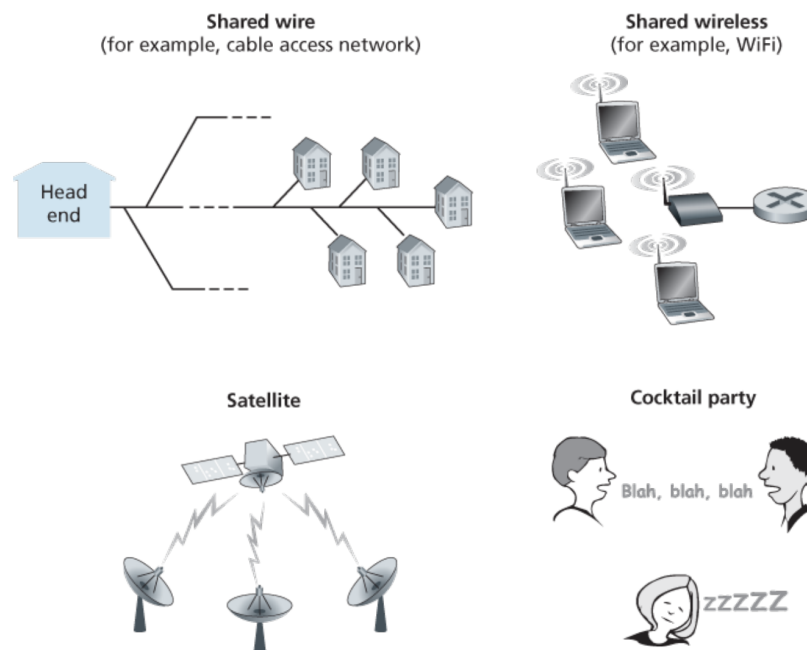


Figure 6.8 Various multiple access channels

- In practice, hundreds or even thousands of nodes can directly communicate over a broadcast channel
- Since all nodes are capable of transmitting frames, more than two nodes can transmit frames at the same time
- When this happens, all receiving nodes receive multiple frames at the same time, such that the transmitted frames *collide* at the receiver
- Typically, when there is a collision, none of the receiving nodes can make sense of any of the frames that were transmitted
- Clearly, if many nodes want to transmit frames frequently, many transmissions will result in collisions, and thus much of the channel's bandwidth will be wasted
- In order to ensure that the broadcast channel performs useful work when multiple nodes are active, we must somehow coordinate the transmissions of the active nodes

- This coordination is the job of the *multiple access protocol*
- Overall, dozens of multiple access protocols have been implemented in a variety of link-layer technologies, however all can generally be classified as belonging to one of three categories
  - *Channel partitioning protocols*
  - *Random access protocols*
  - *Taking-turns protocols*
- Ideally, a multiple access protocol for a broadcast channel of rate  $R$  bits per second should have the following desirable characteristics
  - When *only one* node has data to send, that node has a throughput of  $R$  bps
  - When  $M$  nodes have data to send, each node has a throughput of  $R/M$  bps on average (the instantaneous rates may not be constant)
  - The protocol is decentralized; that is, there is no master node that represents a single point of failure for the network
  - The protocol is simple such that it is inexpensive to implement
- **6.3.1: Channel Partitioning Protocols**
  - We must recall that time-division multiplexing and frequency-division multiplexing are two techniques which can be used to partition a broadcast channel's bandwidth among all nodes sharing that channel

◦




Figure 6.9 A four-node TDM and FDM example

- For example, assume a channel with rate  $R$  bps supports  $N$  nodes
  - Time-division multiplexing divides time into *time frames* and further divides each time frame into  $N$  *time slots*
  - Each slot will be assigned to a node, and when a node has a packet to send, it will do so during its allotted time slot in the TDM time frame
- This scheme is appealing due to the fact that it eliminates the previously mentioned issue of collision and is perfectly fair between all nodes
- Each node gets an average transmission rate of  $R/N$  bps over the course of each time frame
- However, this structure also has two major drawbacks
  - A node is limited to a *maximum* transmission rate of  $R/N$  bps, even if it is the only node with packets to send over the  $R$  bps channel
  - A node also must wait for its allotted time slot, even when, again, it is the only node with any packets to be sent
- While TDM shares the broadcast channel in terms of time, FDM divides the  $R$  bps channel into different frequencies, each with a bandwidth of  $R/N$  and assigns each frequency to one of the  $N$  nodes
- FDM shares both the advantages and drawbacks with TDM
  - It avoids collisions and divides the bandwidth fairly among the  $N$  nodes
  - A node is limited to a *maximum* bandwidth of  $R/N$ , even when it is the only node with packets to be sent
- A third channel partitioning protocol is *code division multiple access (CDMA)*
- CDMA assigns a different *code* to each node, and each node then uses its unique code to encode the data bits it sends
- CDMA also has the unique property that different nodes can transmit simultaneously and yet have their respective receivers correctly receive a sender's encoded data bits

### • 6.3.2: Random Access Protocols

- The second broad class of multiple access protocols are random access protocols

- In a random access protocol, a transmitting node always transmits at the full rate of the channel, namely,  $R$  bps
- If there is a collision, each node involved in the collision will repeatedly retransmit its frame until its frame gets through without a collision
- The node does not immediately retransmit the frame, however, instead waiting a *random delay* before retransmitting the frame
- There are dozens, if not hundreds, random access protocols described in the literature
- We will describe some of the more commonly used ones in this section, starting with *Slotted ALOHA*

- **Slotted ALOHA**

- The slotted ALOHA protocol is one of the simplest random access protocols
- In slotted ALOHA, we will assume the following
  - All frames consist exactly of  $L$  bits
  - Time is divided into slots of size  $L/R$  seconds
    - This means a slot equals the time to transmit one frame
  - Nodes start to transmit frames only at the beginnings of slots
  - The nodes are synchronized so that each node knows when the slot begins
  - If two or more frames collide in a slot, then all the nodes detect the collision event before the slot ends
- We then let  $p$  be a probability (a number between 0 and 1)
  - When the node has a fresh frame to send, it waits until the beginning of the next slot and transmits the entire frame in the slot
  - If there is no collision, the node has successfully transmitted its frame and thus need not consider retransmitting the frame
  - If there *is* a collision, the node detects the collision before the end of the slot, and the node retransmits its frame in each subsequent slot with probability  $p$  until the frame is transmitted without a collision
- The probability  $p$  decides whether the node will retransmit the frame in that slot, or wait until the next slot to try a "biased coin-flip" deciding whether the node will be retransmitted in that frame
- All nodes involved in a collision toss their coins independently
- Unlike channel partitioning, slotted ALOHA allows a node to transmit continuously at the rate  $R$  if that node is the only active node
-



$p$ , or waits for a frame retransmission where it will once again transmit with probability  $p$

- To determine pure ALOHA's maximum efficiency, we will focus on an individual node
- At any given time, the probability that a node is transmitting a frame is  $p$
- Suppose this frame begins transmission at time  $t_0$ 
  - In order for this frame transmission to be successful, no other nodes can begin their transmission in the interval  $[t_0 - 1, t_0]$ , since such a transmission would overlap with the beginning of our node's frame
- By doing a derivation similar to that of slotted ALOHA, we can find that the maximum efficiency of the ALOHA protocol is  $\frac{1}{2e} \approx 0.185$ , which is exactly half that of the slotted ALOHA protocol, which is the major downside when compared to slotted ALOHA

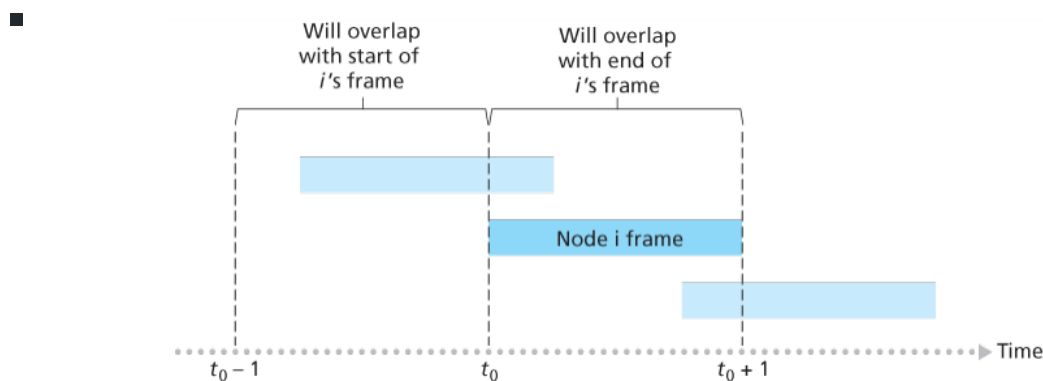



Figure 6.11 Interfering transmissions in pure ALOHA

#### ○ Carrier Sense Multiple Access (CSMA)


- In both of the previous protocols, a node's decision to transmit is made independently of the activity of the other nodes attached to the broadcast channel
- In the CSMA protocols, there are two important rules to limit collisions
  - *Carrier sensing*, such that a node will listen to the channel before transmitting and wait if a frame is being transmitted by another node
  - *Collision detection*, such that a transmitting node listens to the channel while it is transmitting, and if it detects another node sending an interfering frame, it will wait a random time interval and retransmit
- With the property of carrier sensing in place, we then might ask how collisions can even occur in the first place
- This can be best illustrated using a space-time diagram

■



- At time  $t_0$ , node  $B$  senses the channel is idle and begins to transmit, with its bits propagating in both directions along the broadcast medium
- At time  $t_1$ , where  $t_1 > t_0$ , node  $D$  has a frame to send, and node  $B$ 's message has yet to reach  $D$  and thus  $D$  believes the channel is open for its transmission
- In this example, no collision detection occurs
- Let us instead consider the example where collision detection is implemented in the protocol

■



- As we can see from the lower amount of interfering transmissions, this can assist in our channel's efficiency by reducing the amount of useless data sent over the channel
- Let us now summarize the operation of the CSMA with collision detection protocol from the perspective of an adapter in a node attached to a broadcast channel
  - The adapter obtains a datagram from the network layer, prepares a link-layer frame, and puts the frame adapter buffer
  - If the adapter senses that the channel is idle, it starts to transmit the frame, otherwise the adapter will wait until it senses no signal energy and then begin transmission
  - While transmitting, the adapter monitors for the presence of signal energy coming from other adapters using the broadcast channel
  - If the adapter transmits the entire frame without detecting signal energy from other adapters, the adapter is finished with the frame
  - Otherwise, the adapter will abort the transmission, wait a random amount of time, and sense again whether or not the channel is idle
- To calculate the efficiency of CSMA, we will let
  - $d_{prop}$  be the maximum time it takes signal energy to propagate between any two adapters
  - $d_{trans}$  be the time to transmit a maximum size frame
- Then, we can approximate the efficiency,  $E$ , as
  - $$E = \frac{1}{1 + 5\alpha}$$
  - where,
    - $$\alpha = \frac{d_{prop}}{d_{trans}}$$

### • 6.3.3: Taking-Turns Protocols

- There are many of these protocols, and each has many variations, but we will discuss two of the more important ones



- The first is the *polling protocol*, which requires one of the nodes in the broadcast channel to be designated as the *master node*
  - The master node then *polls* each of the nodes in a round-robin fashion
    - This functions by telling node 1 it can transmit up to some maximum number of frames, and then telling node 2 the same once that amount is reached or it senses no signal energy on the channel
  - One drawback of this protocol is that it introduces a *polling delay* since it requires time to notify a node that it can transmit
  - The second taking-turns protocol is the *token-passing protocol*, which does not require the designation of a master node
  - Instead, a small, special-purpose frame, known as the *token* is exchanged among the nodes in some fixed order, achieving a similar effect as the polling protocol
-